


a

Training reward


Mean over the preceding 25 training steps



b

Policy loss

Mean over the preceding 25 training steps



—●— RP-GRPO -△- GRPO -▽- Paired GRPO ···□··· LOO ablation

Training reward is not validation accuracy. Policy-loss scales differ across objectives. Values are the recorded 25-step window means; no additional smoothing is applied.